

EE1007E : Introduction to Digital Design and ICs

Lecture 2 – Design Process

At HotChips'19 Cerebras announced the largest chip in the world at 8.5 in x 8.5in with 1.2 trillion transistors, and 15kW of power, aimed for training of deep-learning neural networks

At HotChips'21, '22 and '23 they showed the next version in 7nm CMOS, with >2x transistor count

- 46,225 mm² silicon
- 2.6 Trillion transistors
- 850,000 AI optimized cores
- 40 Gigabytes on-chip memory
- 20 Petabyte/s memory bandwidth
- 220 Petabit/s fabric bandwidth
- 7nm Process technology at TSMC




Sean Lie,
HotChips'22

Review

- Moore's law is slowing down
 - There are continued improvements in technology, but at a slower pace
- Dennard's scaling has ended two decades ago
 - All designs are now power limited
- Specialization and customization provides added performance
 - Under power constraints and stagnant technology
- Design costs are high
 - Methodology and better reuse to rescue!
 - Abstraction, modularity, regularity are the keys
 - And creativity!

Putting it in Perspective



Performance gains
over the past decade


Lisa Su, HotChips'19 keynote



Digital Logic


Implementing Digital Systems

- Digital systems implement a set of Boolean equations



- How do we actually implement a complex digital system?

Modern (Mostly) Digital Systems-On-A-Chip



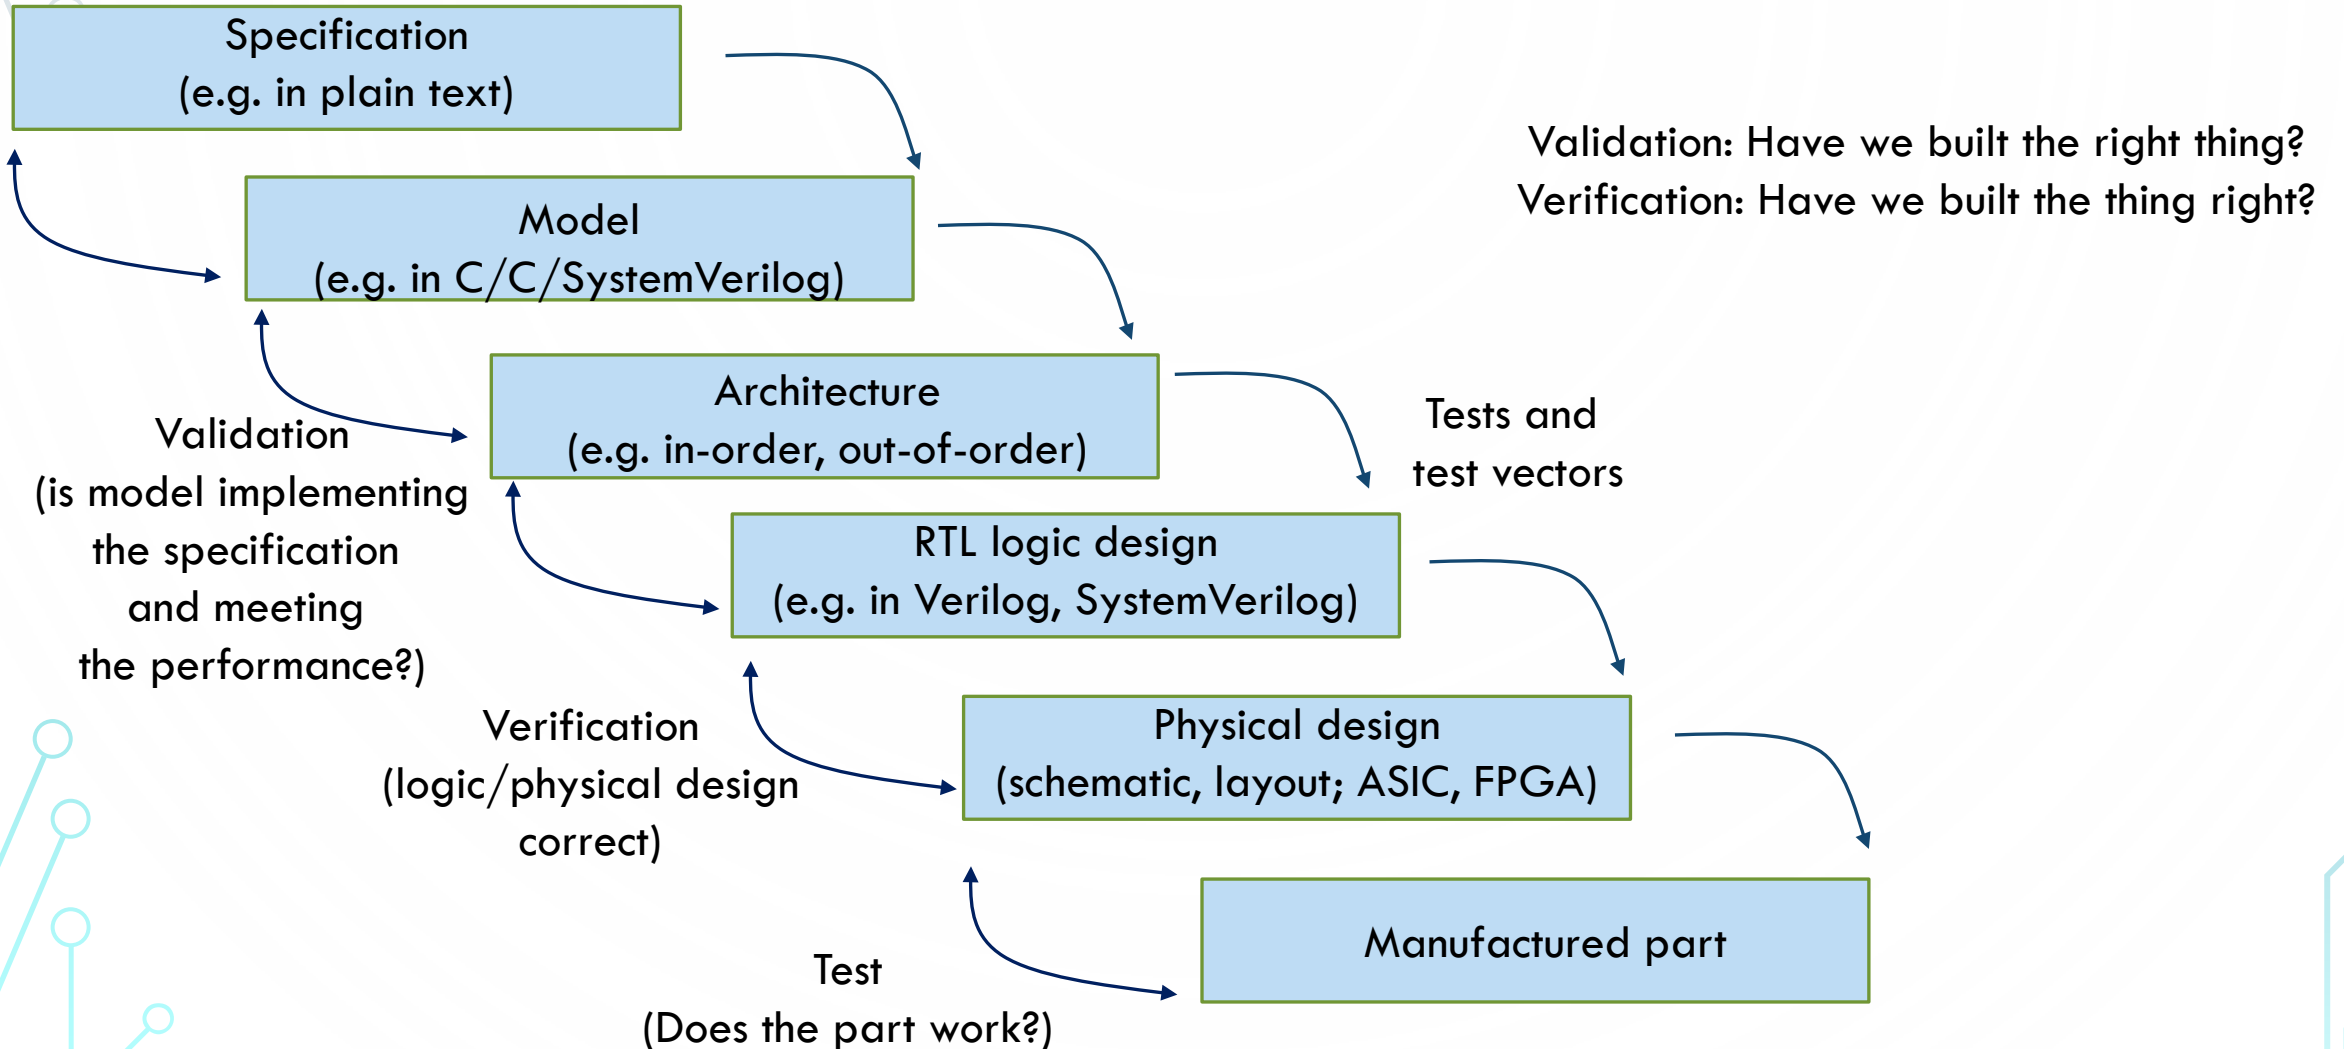
- Multiple large CPUs
- Multiple small CPUs
- GPU
- Neural processing unit (NPU)
- Lots of memory
- DDR memory interfaces

<https://www.semianalysis.com/p/apple-m2-die-shot-and-architecture>

- TSMC N5 (5nm-class) CMOS

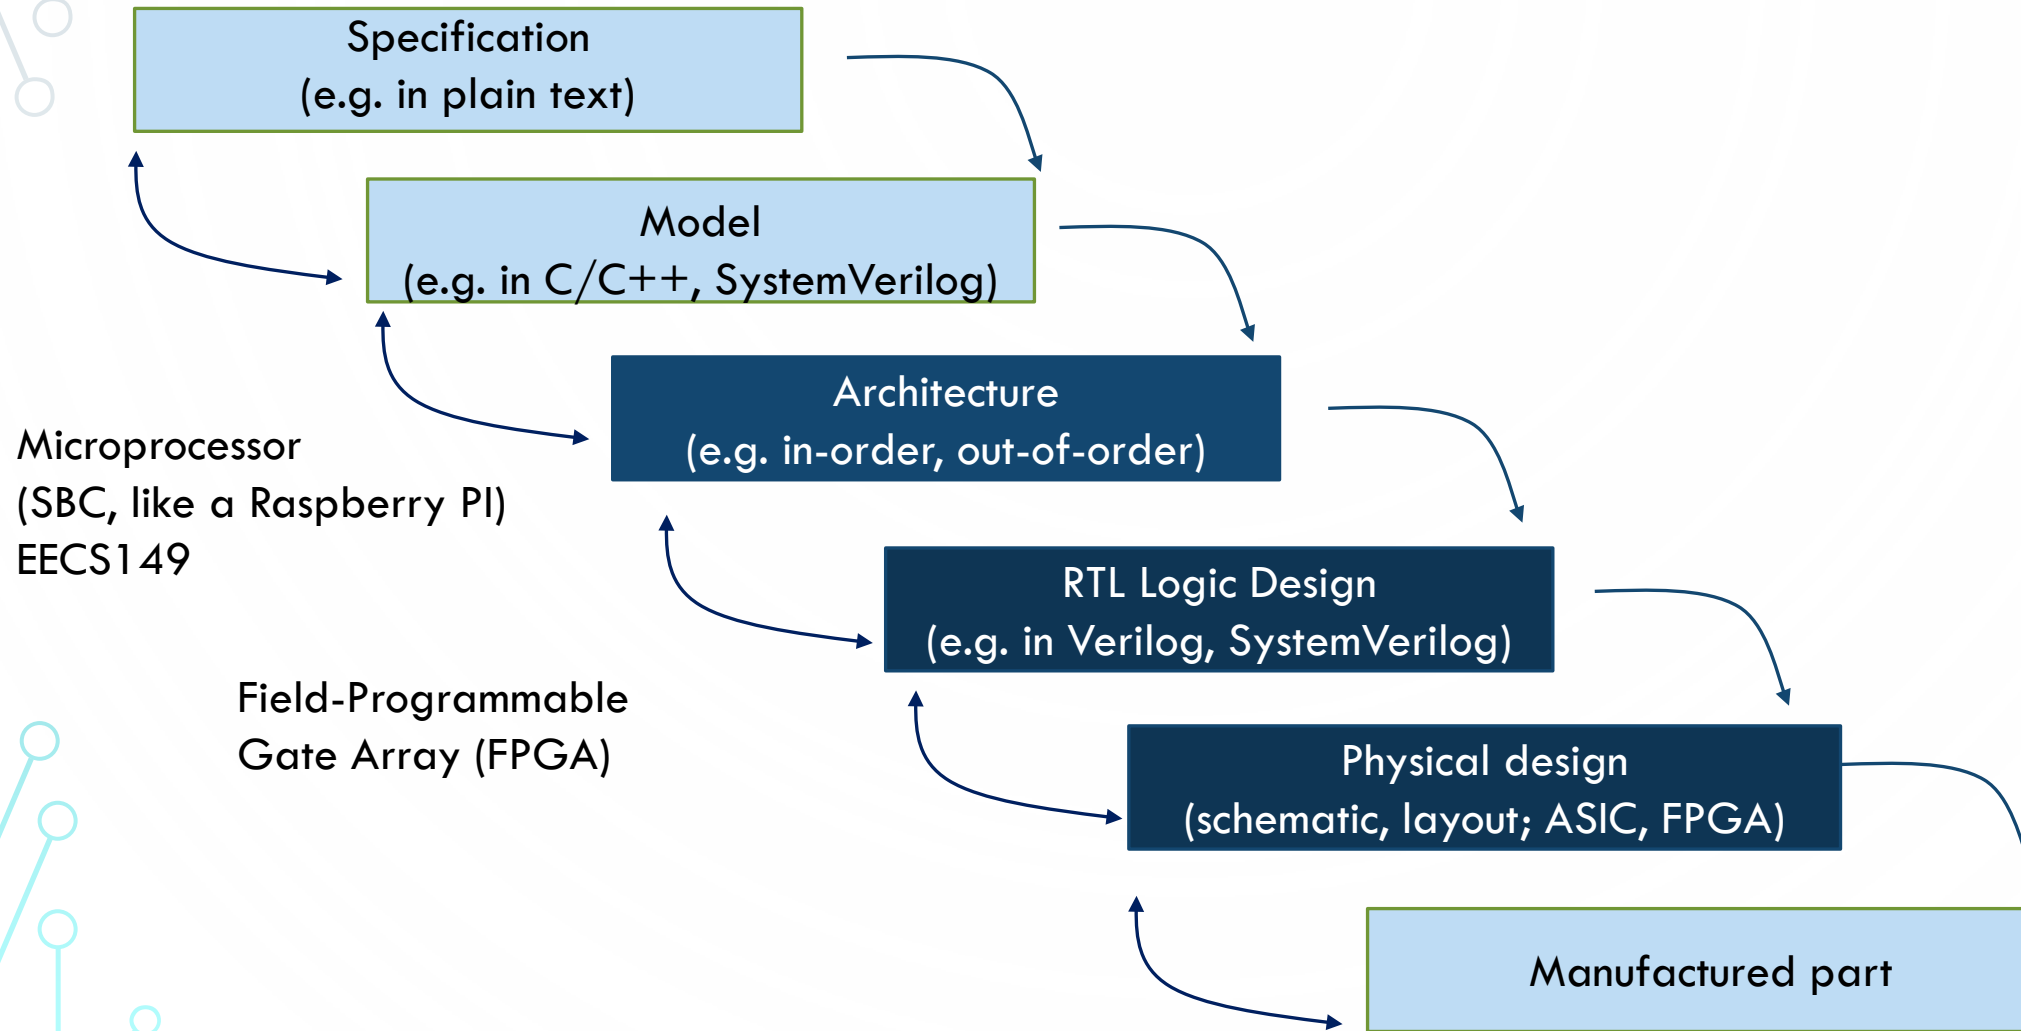
Design Process

- Design through layers of abstractions



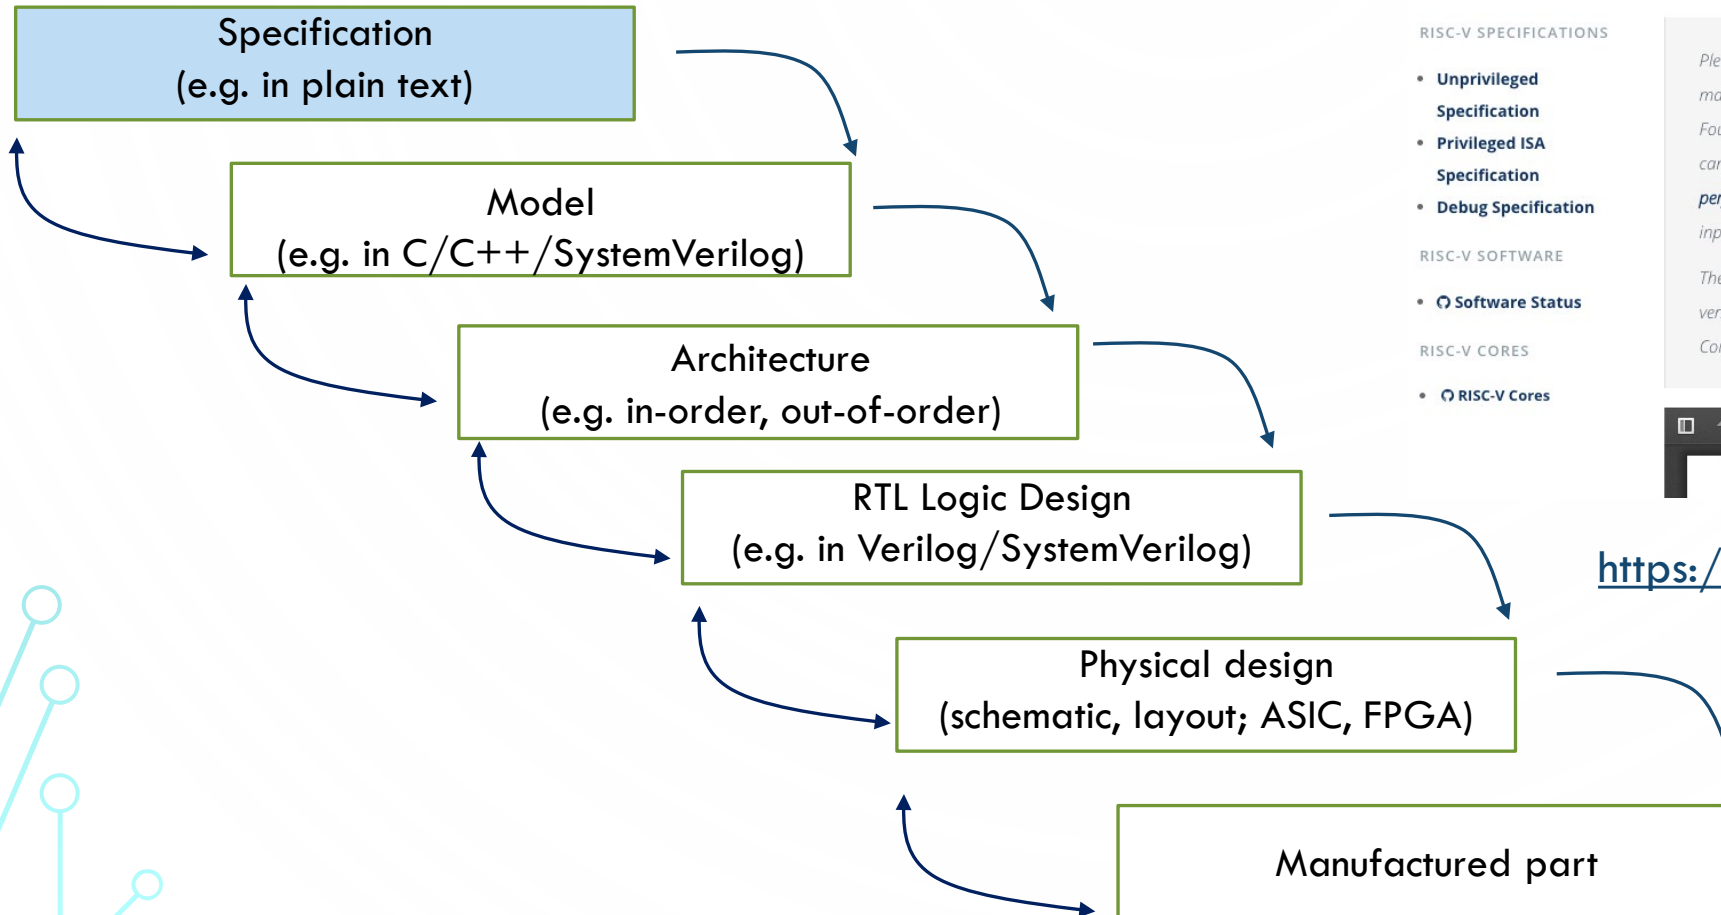
Design Abstractions in EECS151/251A

- Design through layers of abstractions



Example: RISC-V Design Process

- Design through layers of abstractions



ABOUT MEMBERSHIP **SPECS & SUPPORT** CORES & TOOLS NEWS EVENTS

Specifications

Home / Specifications

RISC-V SPECIFICATIONS

- Unprivileged Specification
- Privileged ISA Specification
- Debug Specification

RISC-V SOFTWARE

- Software Status

RISC-V CORES

- RISC-V Cores

Please note, RISC-V ISA and related specifications are developed, ratified and maintained by RISC-V Foundation contributing members within the RISC-V Foundation Technical Committee. Operating details of the Technical Committee can be found in the *RISC-V Foundation Workspace*. Work on the specification is performed on GitHub and the GitHub issue mechanism can be used to provide input into the specification.

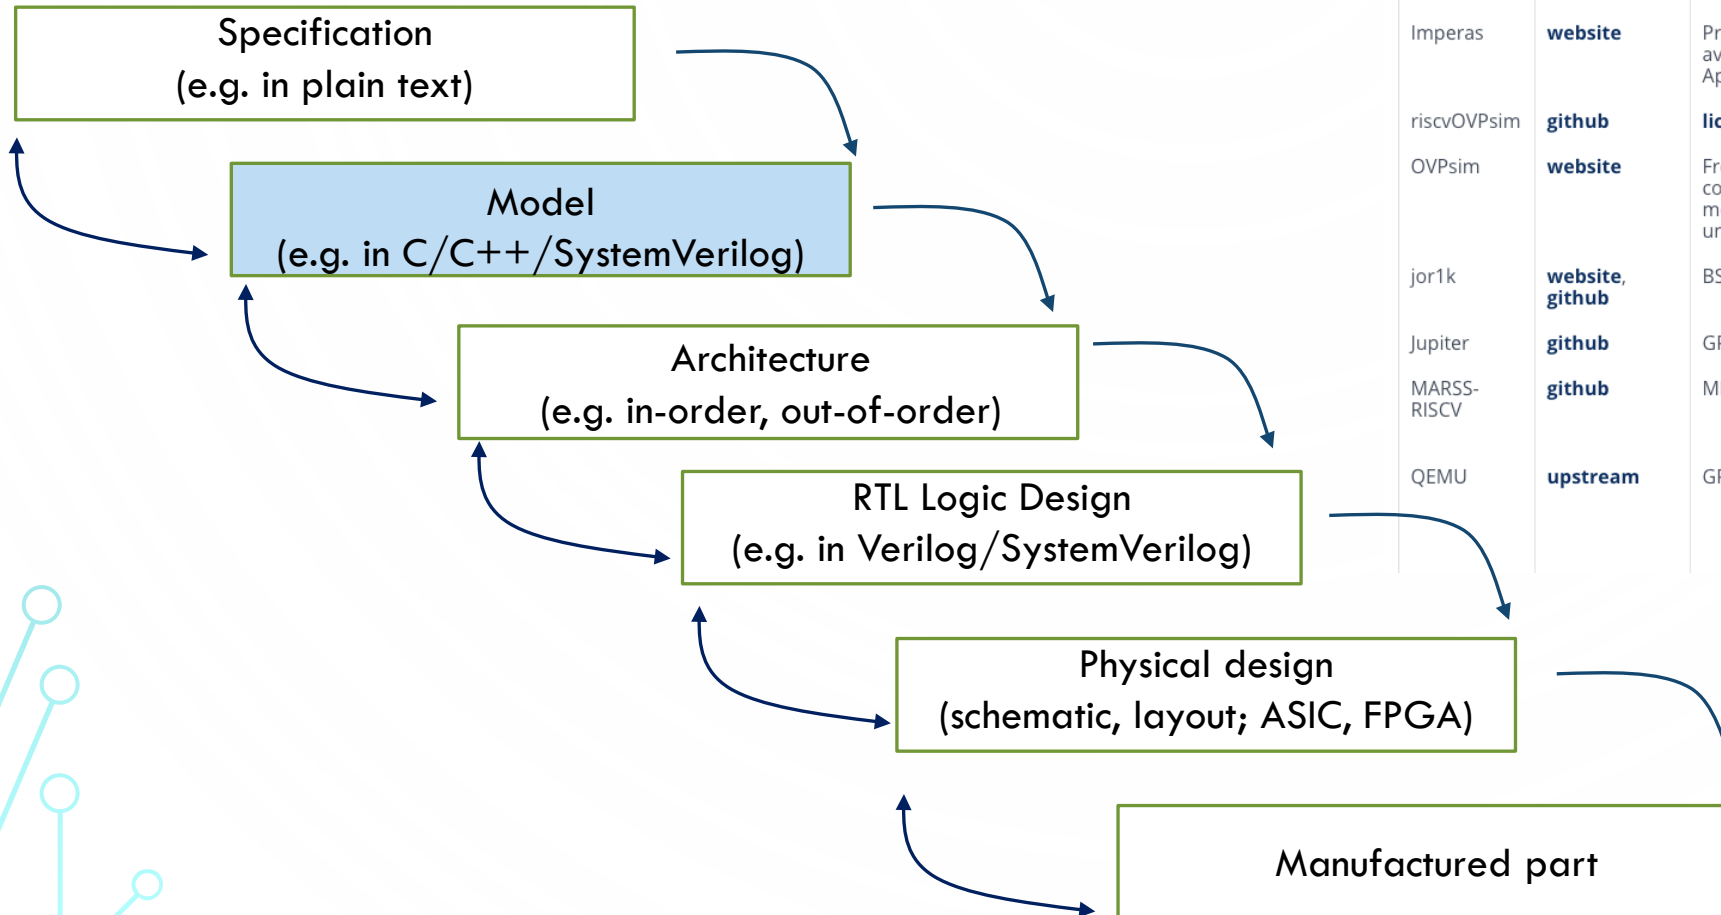
The specifications shown below is the current ratified release. The most recent version of the draft specification, which is in development within the Technical Committee, can be found here on [GitHub](#).



<https://riscv.org/specifications/>

Example: RISC-V Design Process

- Design through layers of abstractions



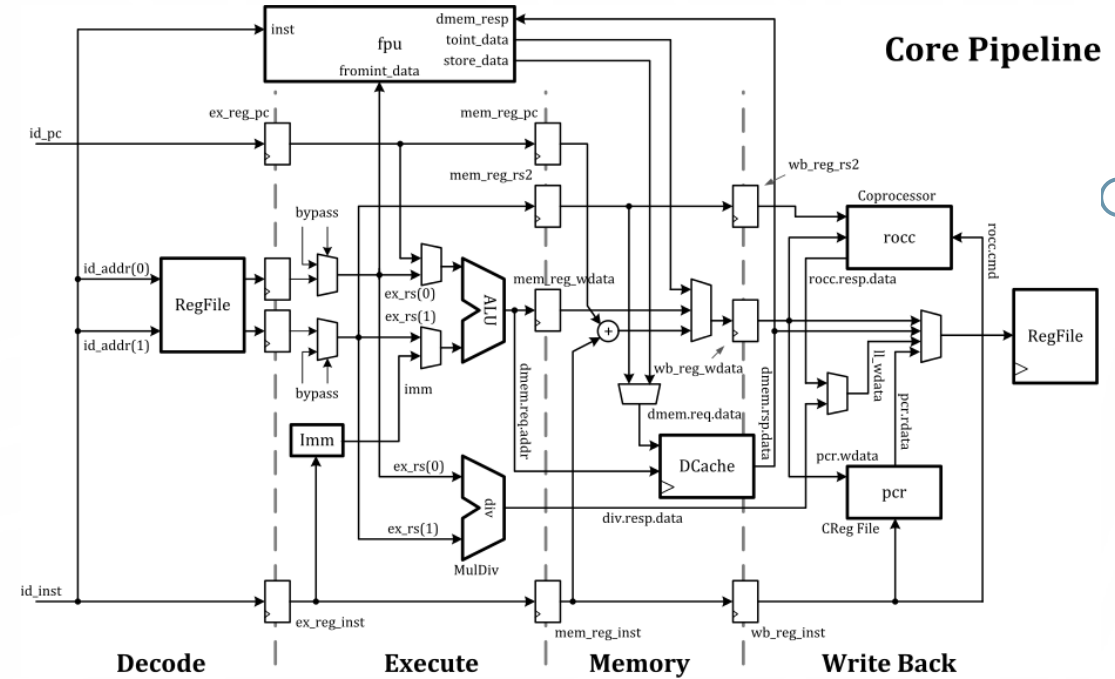
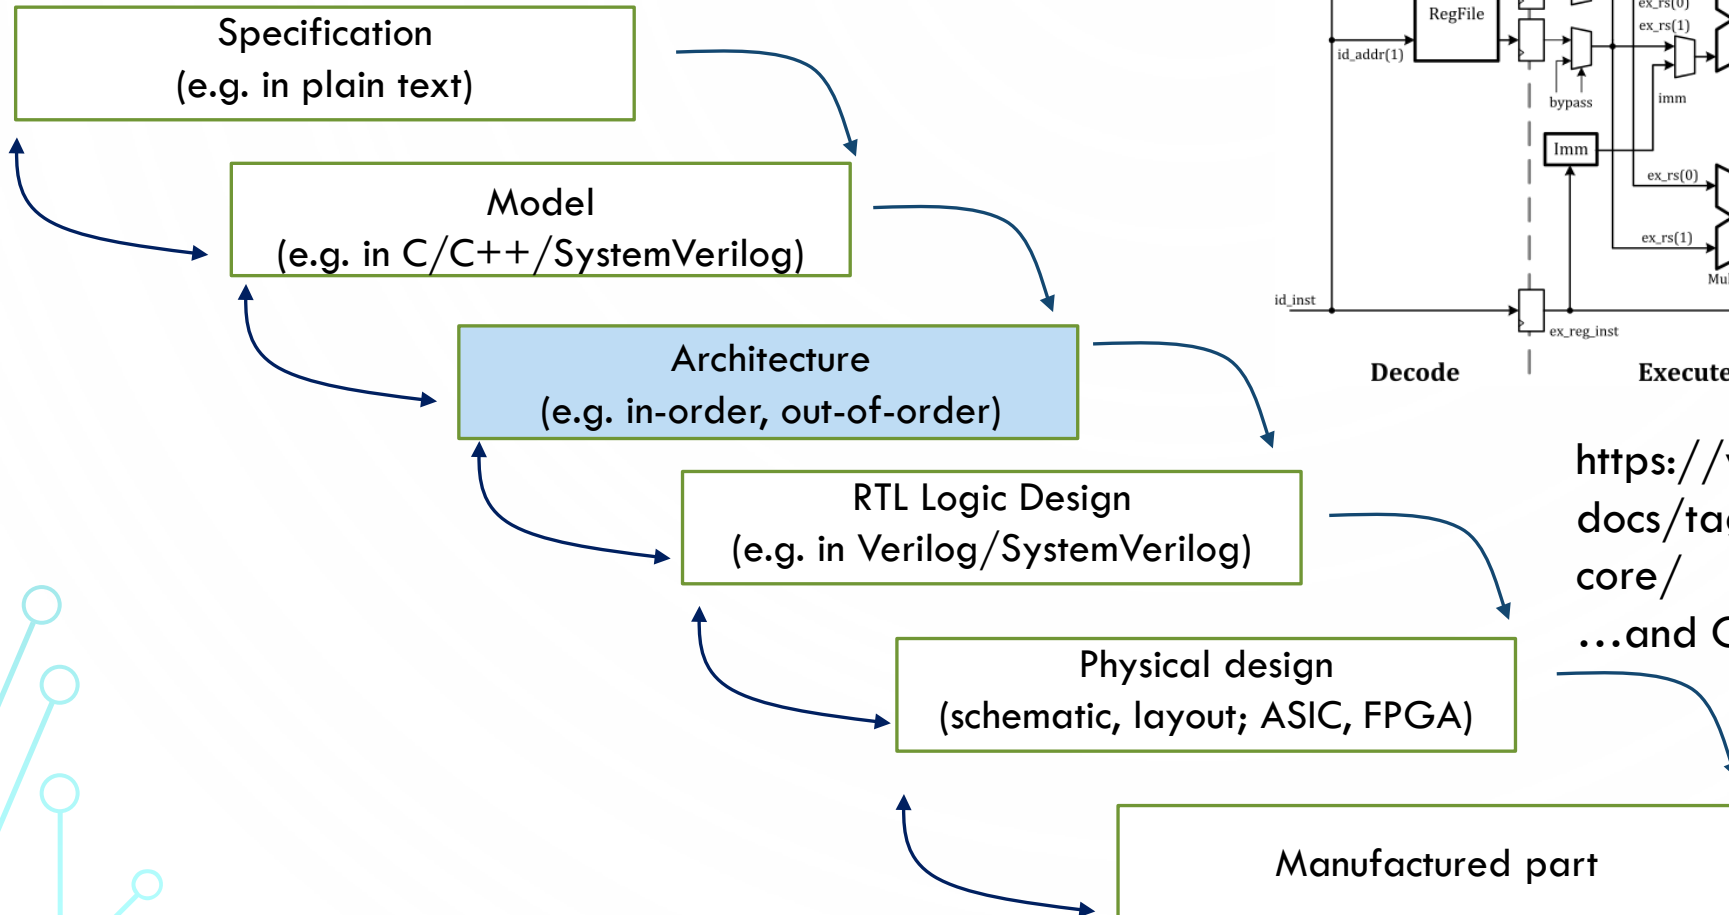
Simulators

Name	Links	License	Maintainers
DBT-RISE-RISCV	github	BSD-3-Clause	MINRES Technologies
FireSim	website , mailing list , github , ISCA 2018 Paper	BSD	Sagar Karandikar, Howard Mao, Donggyu Kim, David Biancolin, Alon Amid, Berkeley Architecture Research
gem5	SW-dev thread , repository	BSD-style	Alec Roelke (University of Virginia)
Imperas	website	Proprietary, models available under Apache 2.0	Imperas
riscvOVPsim	github	license	Imperas
OVPsim	website	Free for non commercial use, models available under Apache 2.0	Imperas
jor1k	website , github	BSD 2-Clause	Sebastian Macke
Jupiter	github	GPL-3.0	Andrés Castellanos
MARSS-RISCV	github	MIT	Gaurav N Kothari, Parikshit P Sarnaik, Gokturk Yuksek (State University of New York at Binghamton)
QEMU	upstream	GPL	Sagar Karandikar (University of California, Berkeley), Bastian Koppelman (University of Paderborn), Alex Suykov, Stefan O'Rear and Michael Clark (SiFive)

<https://riscv.org/software-status/#simulators>

Example: RISC-V Design Process

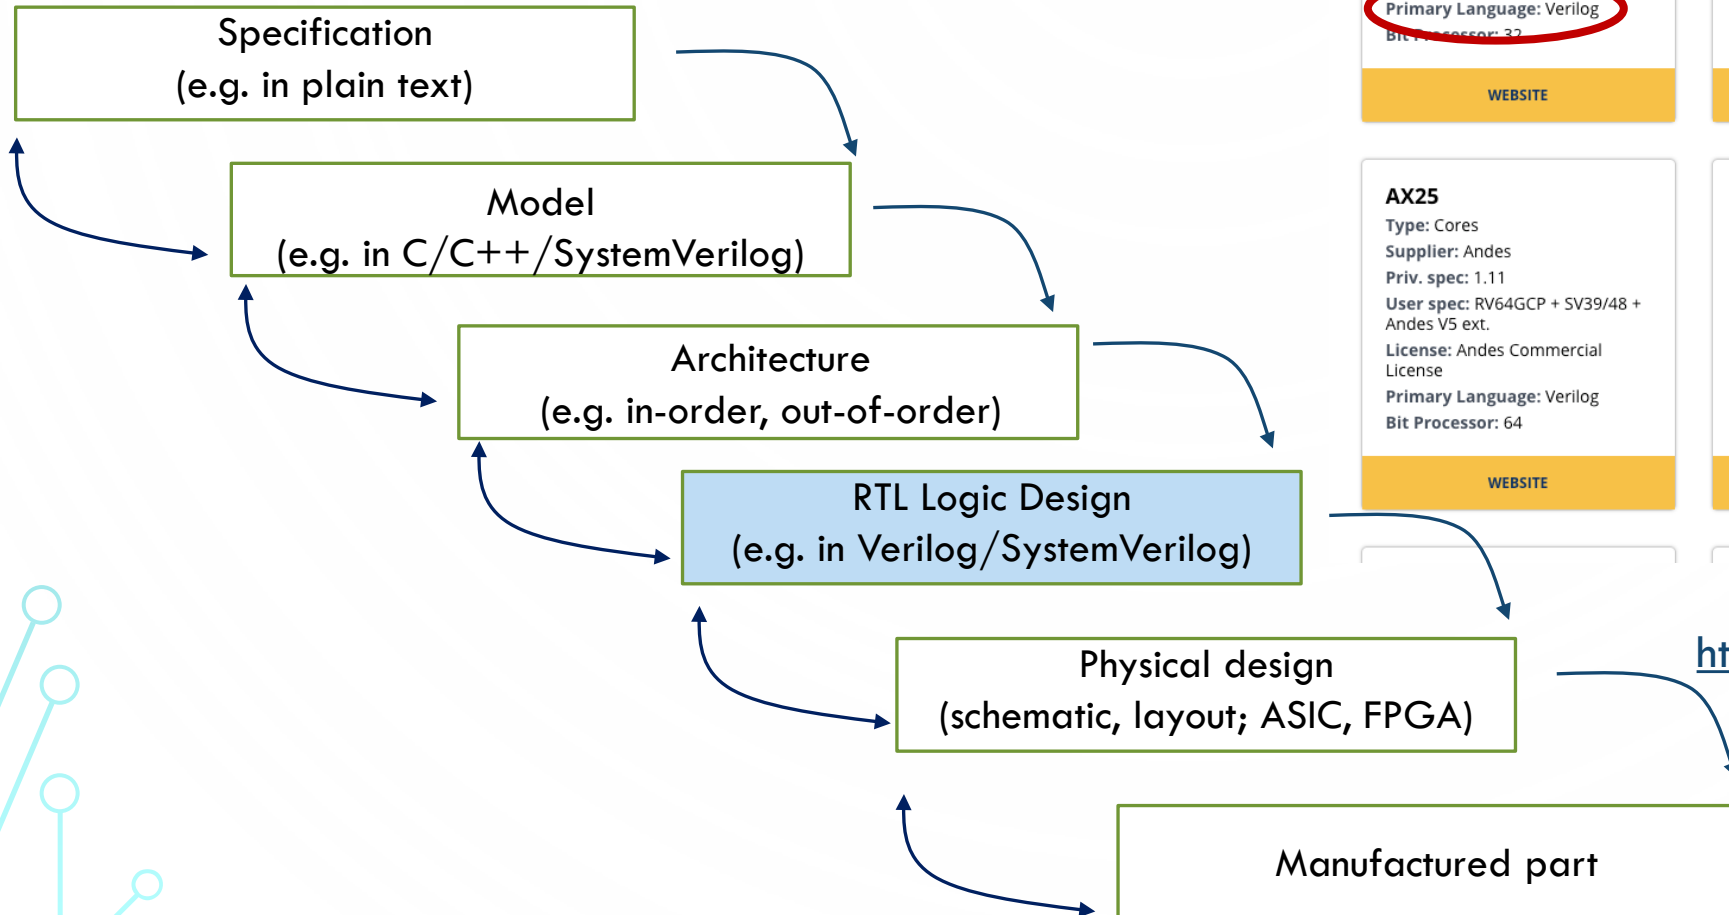
- Design through layers of abstractions



<https://www.lowrisc.org/docs/tagged-memory-v0.1/rocket-core/>
...and CS152

Example: RISC-V Design Process

- Design through layers of abstractions


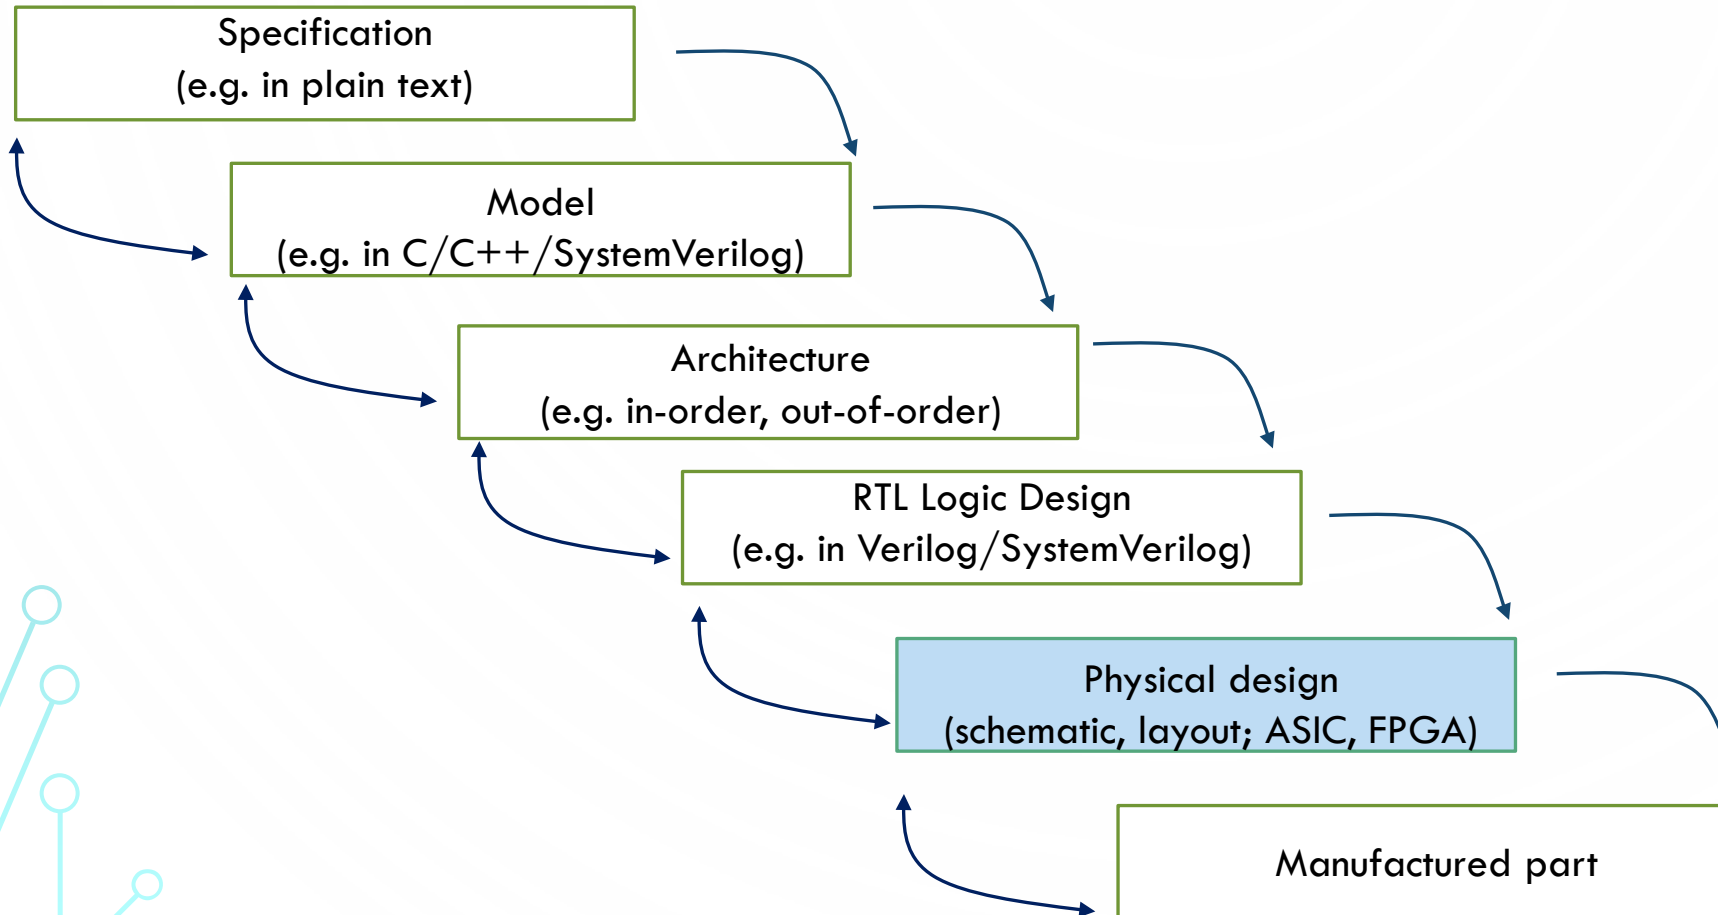


A25 Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV32GCP + SV32 + Andes V5 ext. License: Andes Commercial License Primary Language: Verilog Bit Processor: 32 WEBSITE	A25MP Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV32GCP + SV32 + Andes V5 ext. + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32 WEBSITE	Ariane Type: Cores Supplier: ETH Zurich, Università di Bologna Priv. spec: 1.11-draft User spec: RV64GC License: Solderpad Hardware License v. 0.51 Primary Language: SystemVerilog Bit Processor: 64 WEBSITE GITHUB
AX25 Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV64GCP + SV39/48 + Andes V5 ext. License: Andes Commercial License Primary Language: Verilog Bit Processor: 64 WEBSITE	AX25MP Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV64GCP + SV39/48 + Andes V5 ext. + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 64 WEBSITE	Berkeley Out-of-Order Machine (BOOM) Type: Cores Supplier: Esperanto, UCB Bar Priv. spec: 1.11-draft User spec: 2.3-draft License: BSD Primary Language: Chisel GITHUB

<https://riscv.org/risc-v-cores/>

Example: RISC-V Design Process


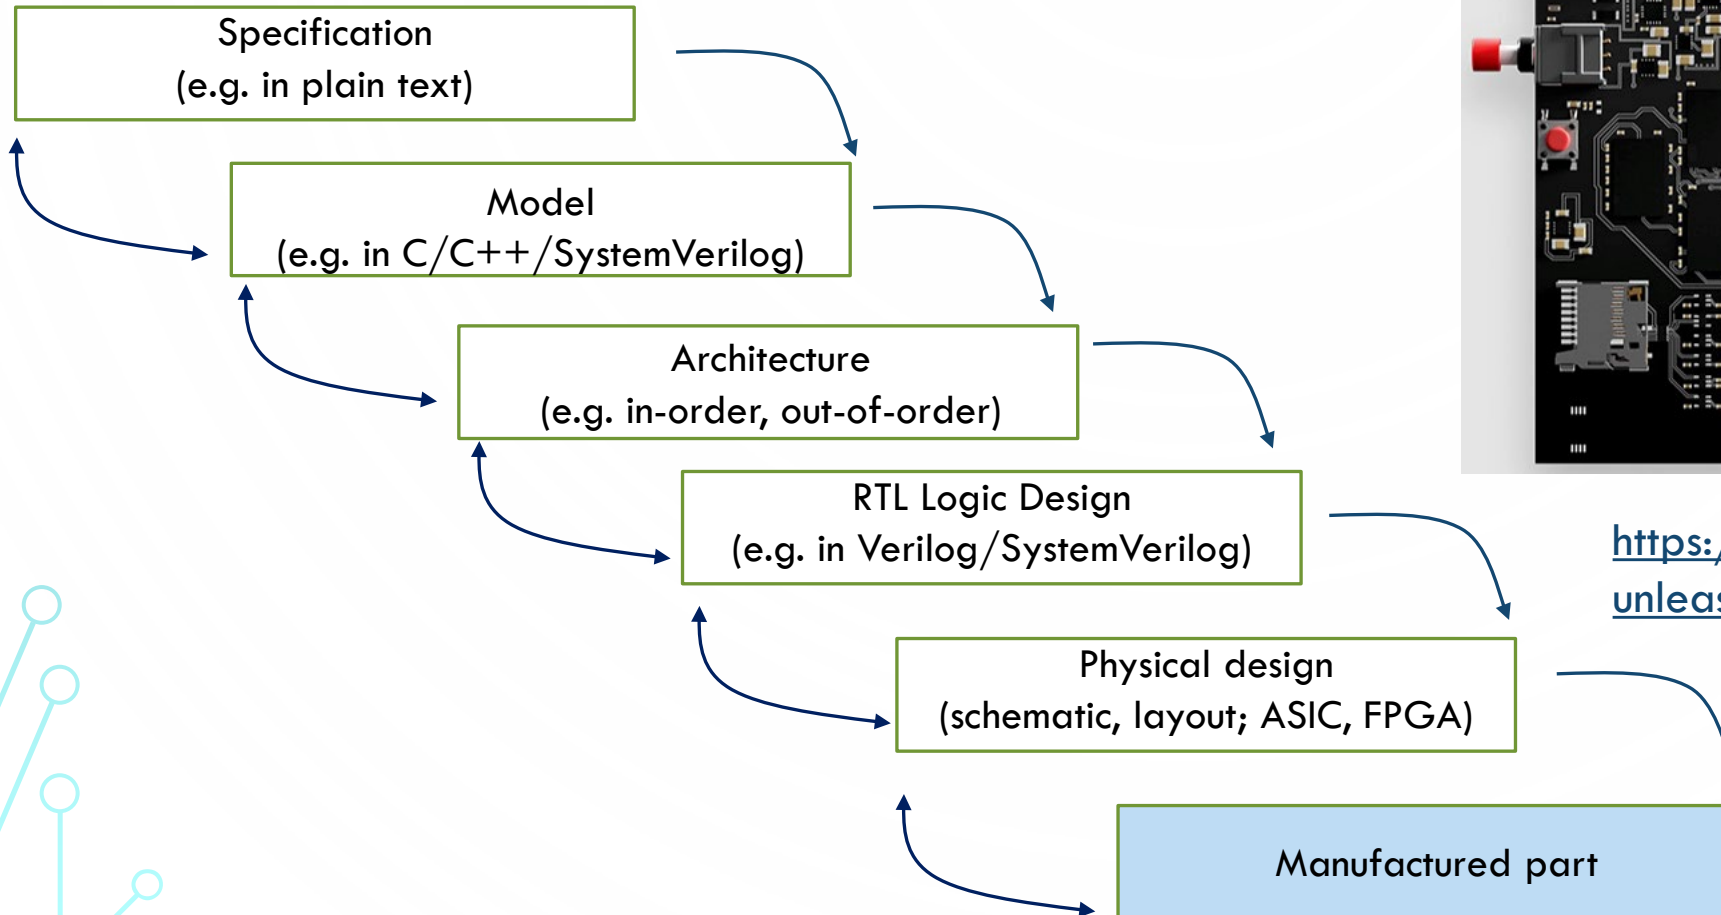
- Design through layers of abstractions



S'22 Class chip
(HotChips'23)

Example: RISC-V Design Process

- Design through layers of abstractions



<https://www.sifive.com/boards/hifive-unleashed>

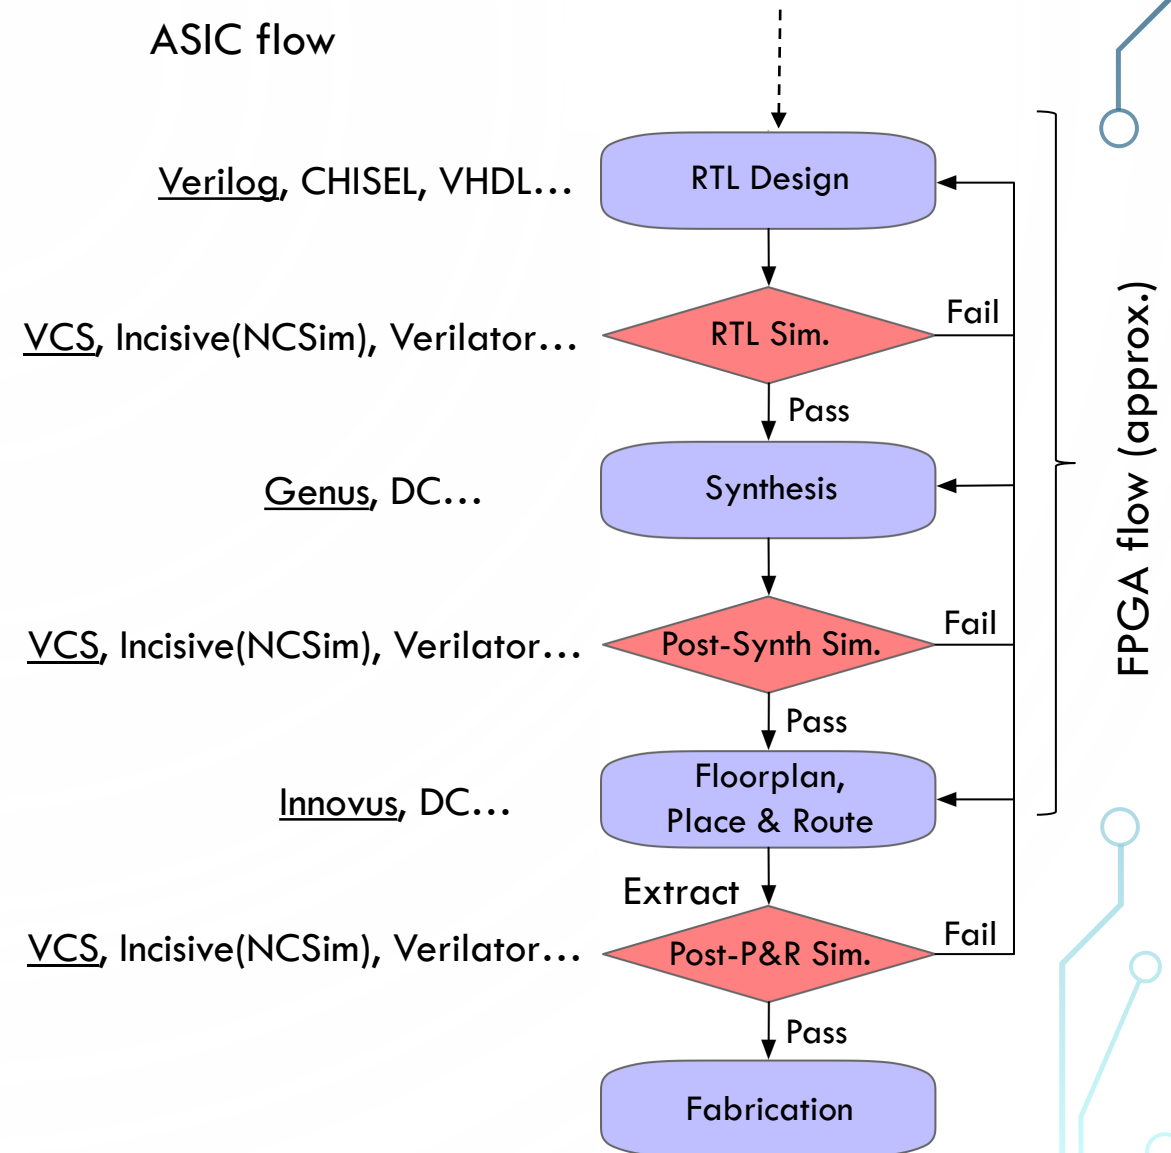
RTL → Physical Design

RTL Logic Design
(e.g. in Verilog/SystemVerilog)

Physical design
(schematic, layout; ASIC, FPGA)

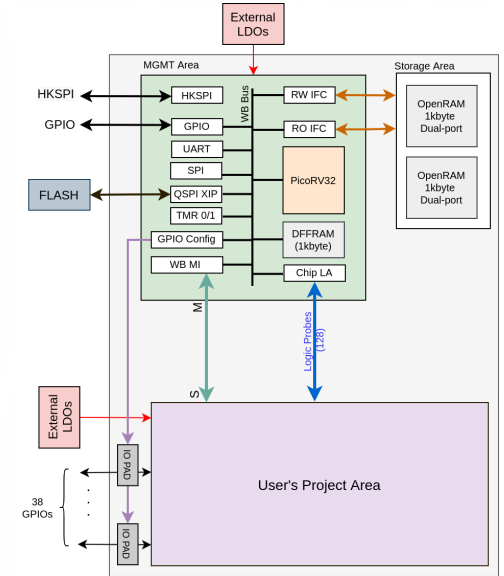
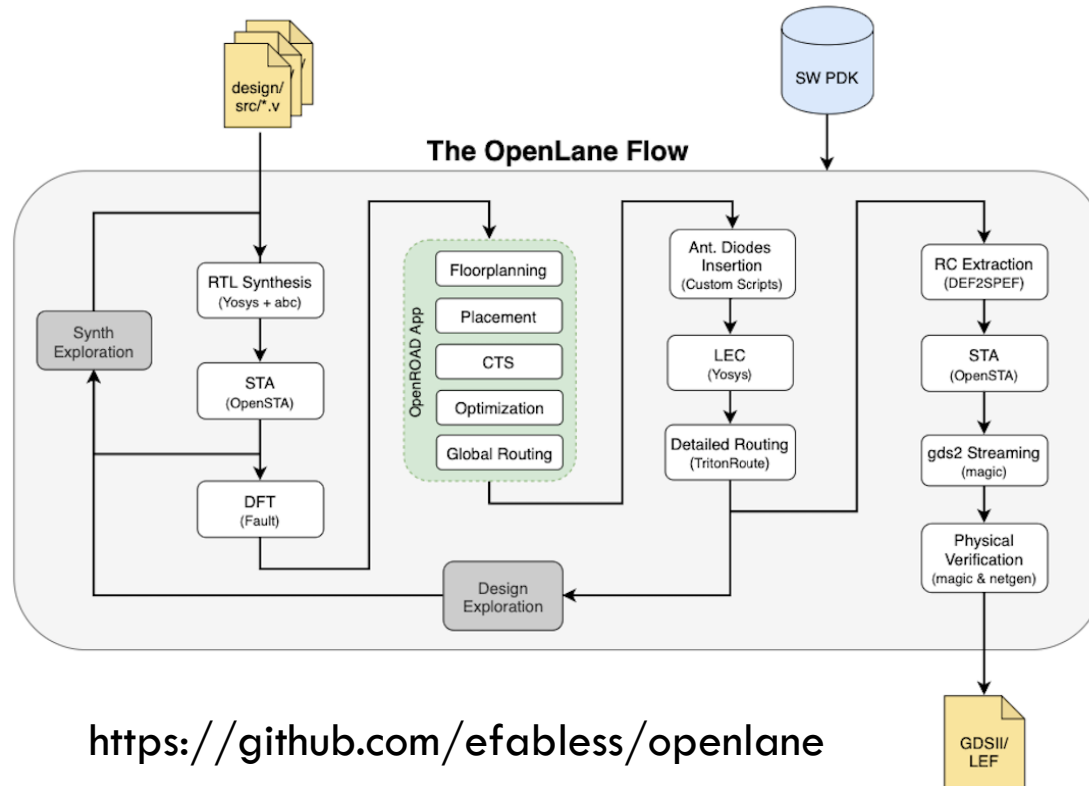
- Labs focus on a process of translating RTL to physical ASIC or FPGA by using industry-standard tools.
- Explores the entire design stack.

ASIC flow




Open-Source Flows

- SkyWater 130nm is an open-source design kit
- OpenROAD (UCSD) and OpenLane (eFabless) are open-source design flows
 - Work with Sky130
 - ASIC labs can target SkyWater130nm with commercial tools



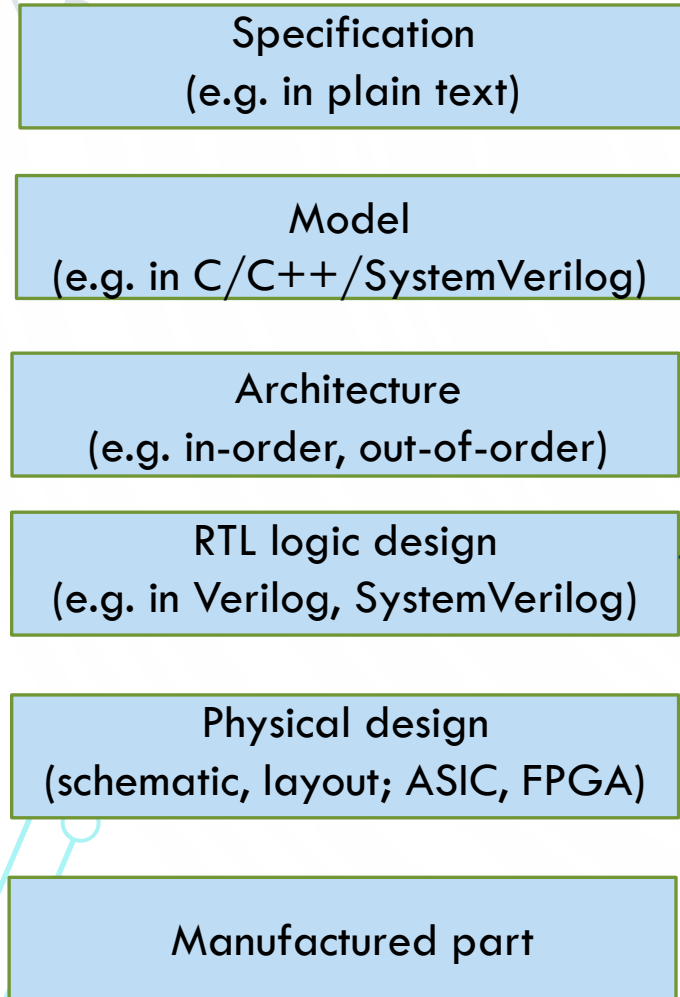
<https://github.com/efabless/caravel>



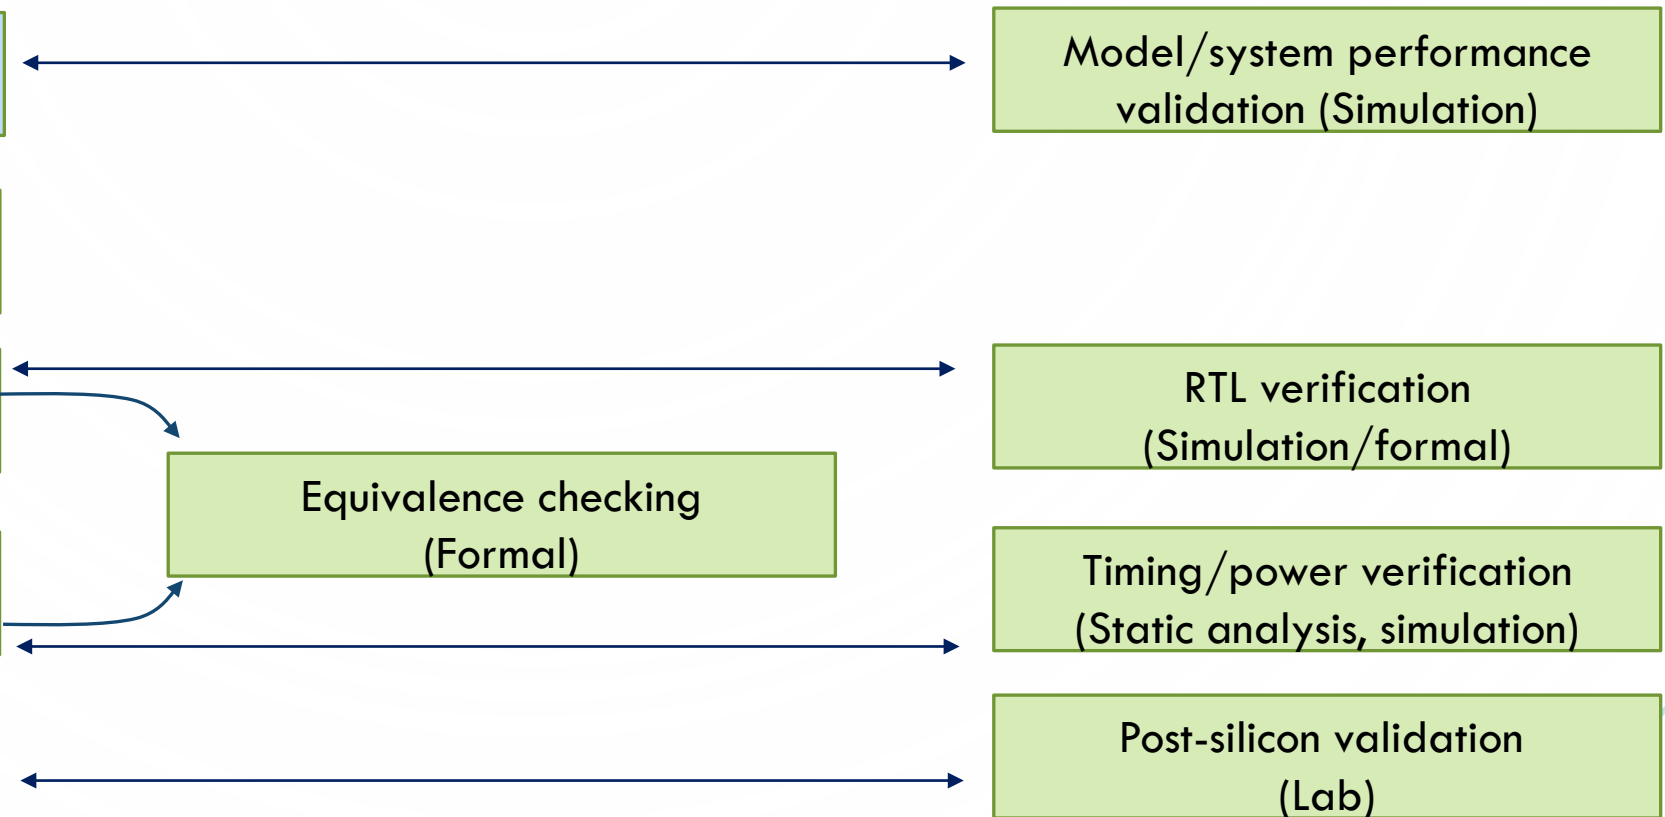
<https://efabless.com/projects/35> 16


Design, Verification and Validation

- Design steps



- Verification and validation steps





Boolean Logic in A Nutshell

Boolean Logic and Logic Gates (From CS61C/EE16B)

- Logic gates

Name

Boolean equation


Symbol

NOT/INV

Truth table

NOT or Inverter

$$\text{Out} = \overline{A}$$




A	Out
0	1
1	0

Single input

Buffer


$$\text{Out} = A$$



A	Out
0	0
1	1

NAND

$$\text{Out} = \overline{A \cdot B}$$




A	B	Out
0	0	1
0	1	1
1	0	1
1	1	0

A	B	Out
0	0	1
0	1	0
1	0	0
1	1	0

NOR

$$\text{Out} = \overline{A + B}$$



- In CMOS, basic logic gates are inverting

More Logic Gates

Name


Boolean equation

Symbol

Truth table


AND

$$\text{Out} = A \cdot B$$




A	B	Out
0	0	0
0	1	0
1	0	0
1	1	1

In CMOS




OR

$$\text{Out} = A + B$$



A	B	Out
0	0	0
0	1	1
1	0	1
1	1	1

In CMOS



More Logic Gates

Name


Boolean equation

Symbol

Truth table

Exclusive OR
XOR


$$\text{Out} = A \oplus B$$



A	B	Out
0	0	0
0	1	1
1	0	1
1	1	0

Exclusive NOR
XNOR

$$\text{Out} = \overline{A \oplus B}$$



A	B	Out
0	0	1
0	1	0
1	0	0
1	1	1


- XOR and XNOR are both inverting and non-inverting

Multi-Input Gates

3-Input NAND

NAND3 Boolean equation

$$\text{Out} = \overline{A \cdot B \cdot C}$$




A	B	C	Out
0	0	0	1
0	0	1	1
0	1	0	1
0	1	1	1
1	0	0	1
1	0	1	1
1	1	0	1
1	1	1	0

And-Or-Invert

AOI21 Boolean equation

$$\text{Out} = \overline{A \cdot B + C}$$




A	B	C	Out
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	0

- Single gate in modern CMOS usually doesn't have more than 3-4 inputs

Combinational Logic (CL) Blocks

Example four-input function:




Truth Table

A	B	C	D	Out
0	0	0	0	F(0,0,0,0)
0	0	0	1	F(0,0,0,1)
0	0	1	0	F(0,0,1,0)
0	0	1	1	F(0,0,1,1)
0	1	0	0	F(0,1,0,0)
0	1	0	1	F(0,1,0,1)
0	1	1	0	F(0,1,1,0)
0	1	1	1	F(0,1,1,1)
1	0	0	0	F(1,0,0,0)
1	0	0	1	F(1,0,0,1)
1	0	1	0	F(1,0,1,0)
1	0	1	1	F(1,0,1,1)
1	1	0	0	F(1,1,0,0)
1	1	0	1	F(1,1,0,1)
1	1	1	0	F(1,1,1,0)
1	1	1	1	F(1,1,1,1)

- Output a function only of the current inputs (no history).
- Truth-table representation of function. Output is explicitly specified for each input combination.
- In general, CL blocks have more than one output signal, in which case, the truth-table will have multiple output columns.

Example CL Block

- 2-bit adder. Takes two 2-bit integers and produces 3-bit result.



A1	A0	B1	B0	C2	C1	C0
0	0	0	0	0	0	0
0	0	0	1	0	0	1
0	0	1	0	0	1	0
0	0	1	1	0	1	1
0	1	0	0	0	0	1
0	1	0	1	0	1	0
0	1	1	0	0	1	1
0	1	1	1	1	0	0
1	0	0	0	0	1	0
1	0	0	1	0	1	1
1	0	1	0	1	0	0
1	0	1	1	1	0	1
1	1	0	0	0	1	1
1	1	0	1	1	0	0
1	1	1	0	1	0	1
1	1	1	1	1	1	0

- Think about truth table for 32-bit adder. It's possible to write out, but it might take a while!

Theorem:

Any combinational logic function can be implemented as a network of simple logic gates.

Quiz

Total number of possible truth tables with 4 inputs is:

- a) 4
- b) 16
- c) 256
- d) 16,384
- e) 65,536
- f) None of the above


www.yellkey.com/notice

Logic Circuit

- A logic gate can be implemented in different ways



NAND

$$\text{Out} = \overline{A \cdot B}$$




A	B	Out
0	0	1
0	1	1
1	0	1
1	1	0

CMOS




Sizing of transistors (W/L) in CMOS changes properties (delay, power, size) of a logic gate



Mechanical LEGO logic gates. A clockwise rotation represents a binary “one” while a counter-clockwise rotation represents a binary “zero.”

Sequential Logic Blocks



- Output is a function of both the current inputs and the state.
- State represents the memory.
- State is a function of previous inputs.
- In synchronous digital systems, state is updated on each clock tick.


Flip-Flop as A Sequential Circuit

- Synchronous state element transfers its input to the output on a rising (or, rarely, falling) clock edge



- Flip-flop

- Rising edge

Signifies 'edge triggered'



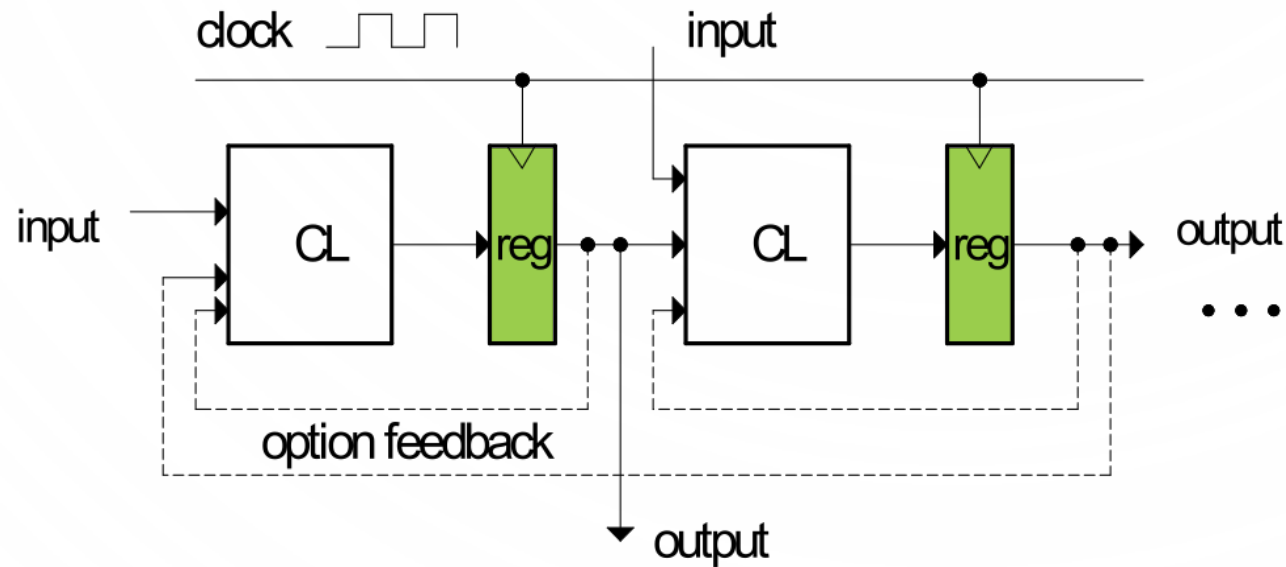
- 4-bit register



Register Transfer Level Abstraction (RTL)

Any synchronous digital circuit can be represented with:

- Combinational Logic (CL) blocks, plus
- State elements (registers or memories)
- Clock orchestrates *sequencing* of CL operations



- State elements are combined with CL blocks to control the flow of data.

Administrivia



- Labs and discussions start this week
- Lab 1 posted, please start it before coming to the lab session
- Lab 2 is more involved
 - Be prepared
 - Verilog primer
- Homework 1 posted this week, due next Friday
 - Start early
- Apple NSI event: Wednesday, August 30, 11:45am, Banatao Auditorium
 - Burritos + bobba



Design Metrics: Robustness



What Makes Circuits Digital?

- Chips are noisy
- Supply noise will appear at the output of the logic gate



- The following logic gate should still interpret its inputs as 0s and 1s
- This necessary property is called "Restoration" or "Regeneration"
- A lot of money was spent in the past to unsuccessfully make logic out of non-regenerative gates
 - Some of emerging CMOS replacements don't have gain...


Beneath the Digital Abstraction




- Logic levels:
 - Mapping a continuous voltage onto a discrete binary logic variable
 - Low (0): $[0, V_L]$
 - High (1): $[V_H, V_{DD}]$
 - V_L V_H : nominal voltage levels

Voltage Transfer Characteristic

- A gate should interpret everything that is close to 0V as a logic 0
 - And everything close to V_{DD} as a logic 1




a) Ideal



b) Real

Nominal Voltage Levels





$$\begin{aligned} V_{OH} &= f(V_{OL}) \\ V_{OL} &= f(V_{OH}) \\ V_M &= f(V_M) \end{aligned}$$

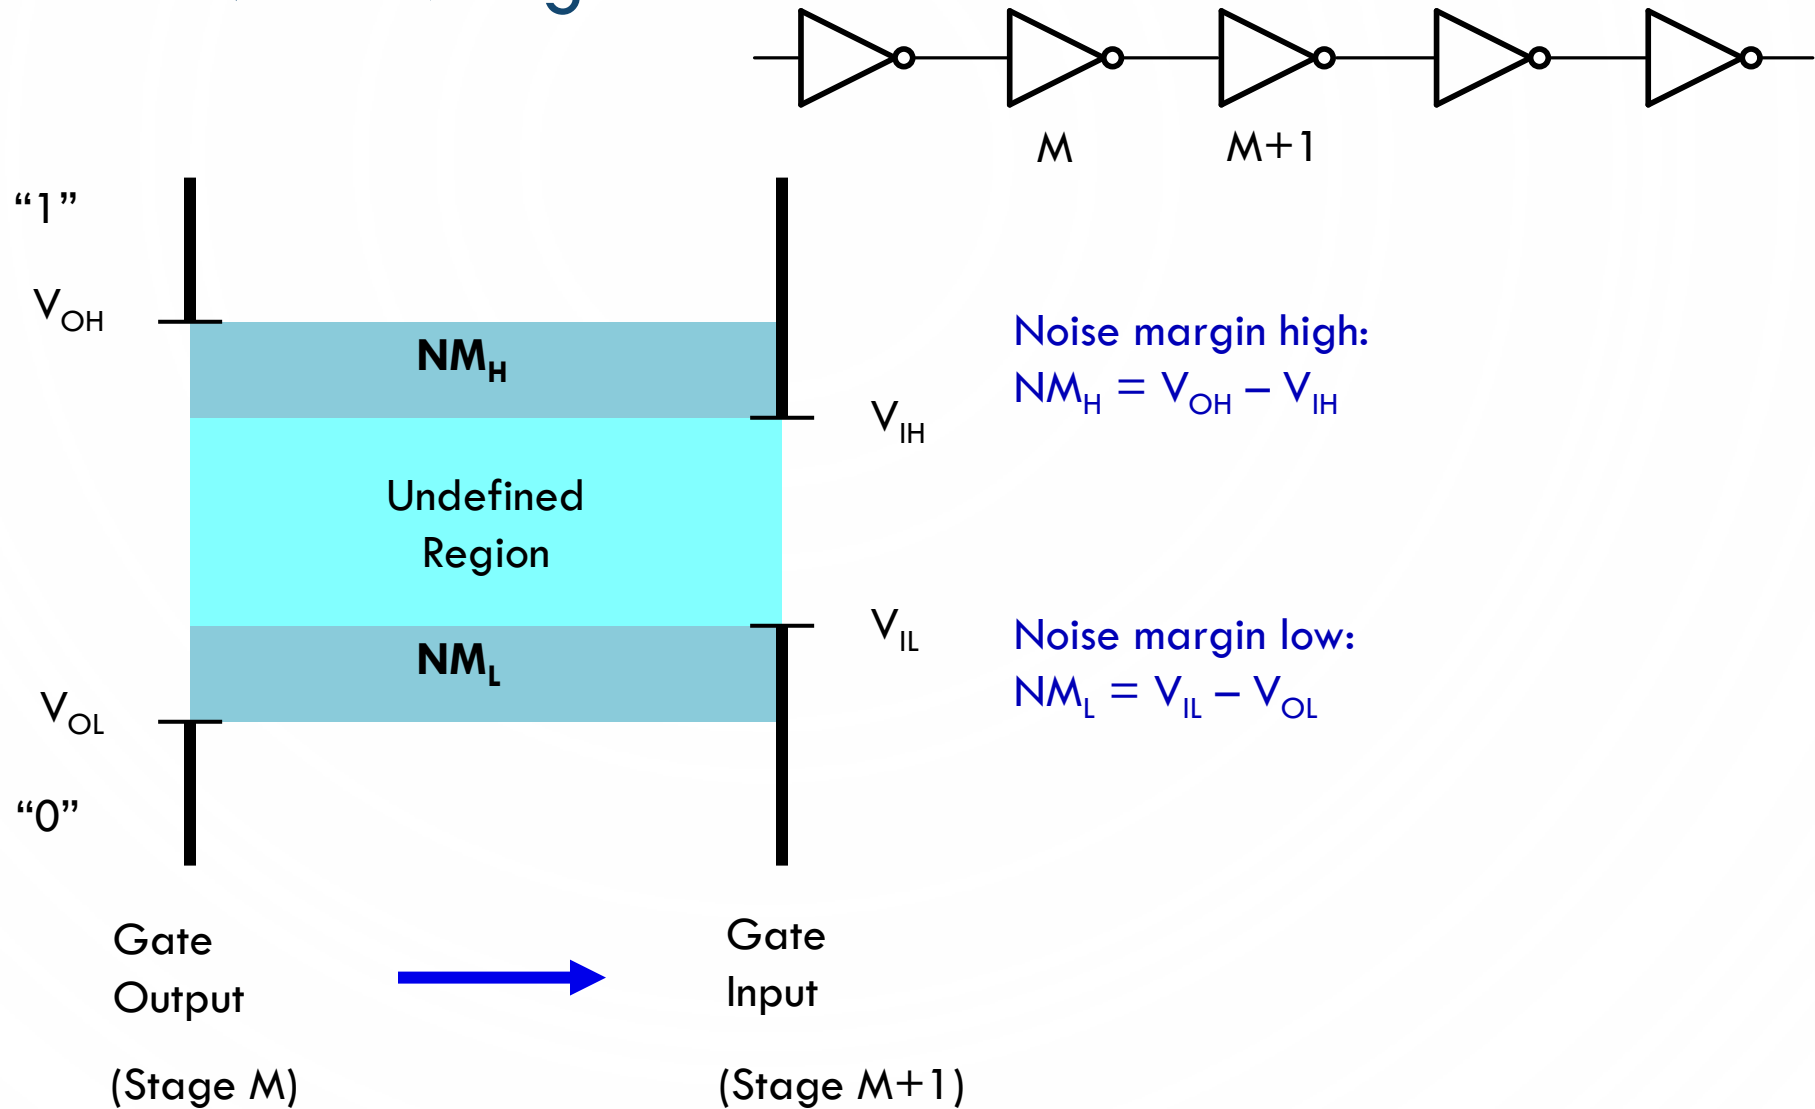
In CMOS:

$$\begin{aligned} V_{OH} &= V_{DD} \\ V_{OL} &= 0 \\ V_M &\sim V_{DD}/2 \end{aligned}$$

Mapping Between Analog Voltages and Digital Signals




Definition of Noise Margins



The amount of **noise** that could be added to a **worst-case output** so that the signal can still be interpreted correctly as a **valid input** to the next gate.

Regenerative Property

- Ensures that a **disturbed** signal gradually **regenerates** one of the **nominal voltage levels** after passing through a few logical stages.
 - Look for a sharp transition in voltage transfer characteristics.






Design Metrics: Performance



Design Tradeoffs

- The desired functionality can be implemented with different performance, power or cost targets




Digital Logic Delay


- Changes at the inputs do not instantaneously appear at the outputs
 - There are finite resistances and capacitances in each gate...




- Propagation through a chain of gates



Delay Definitions

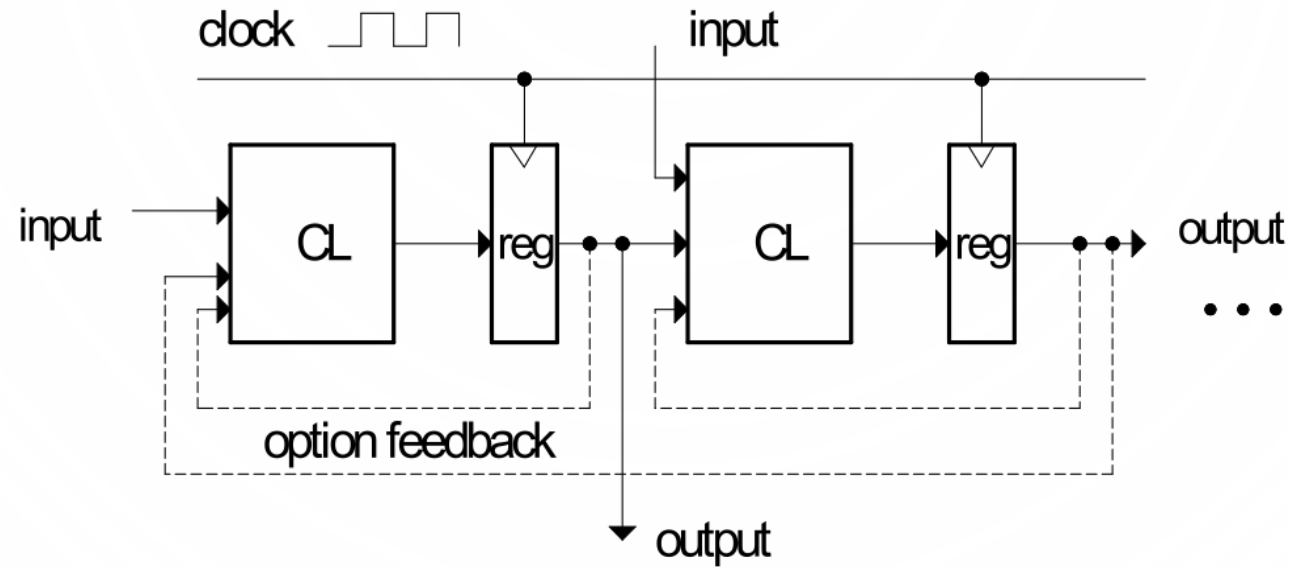


- Delay calculations need to be additive
 - Calculate the delay from the same point in the waveform



Digital Logic Timing

- The longest propagation delay through CL blocks sets the maximum clock frequency



- To increase clock rate:
 - Find the longest path
 - Make it faster



Performance

- Throughput

- Number of tasks performed in a unit of time (operations per second)
- E.g. Google TPUv4 chip performs 275 TFLOPS (10^{12} floating-point operations per second, where a floating point operation is BFLOAT16)
- Watch out for 'op' definitions – can be a 1-b ADD or a double-precision FP add (or more complex task)
- Peak vs. average throughput

- Latency

- How long does a task take from start to finish
- E.g. facial recognition on a phone takes 10's of ms
- Sometime expressed in terms of clock cycles
- Average vs. 'tail' latency






Design Metrics: Energy and Power

Energy and Power

- Energy (in joules (J))
 - Needed to perform a task
 - Add two numbers or fetch a datum from memory
 - (or fetch two numbers, add them and store in memory)
 - Active and standby
 - Battery stores certain amount of energy (in $Ws = J$ or Wh)
 - That is what utility charges for (in kWh)
 - Energy footprint matters for the environment!
- Power (in watts (W))
 - Energy dissipated in time ($W = J/s$)
 - Sets cooling requirements
 - Heat spreader, size of a heat sink, forced air, liquid, ...

TPU v4. ~200W





Design Metrics: Cost


Cost

- Non-recurring engineering (NRE) costs
- Cost to develop a design (product)
 - Amortized over all units shipped
 - E.g. \$20M in development adds \$.20 to each of 100M units
- Recurring costs
 - Cost to manufacture, test and package a unit
 - Processed wafer cost is ~10k (around 16nm node) which yields:
 - 1 Cerebras wafer-scale chip
 - 50-100 large FPGAs or GPUs
 - 500 laptop CPUs
 - >1000 cell phone SoCs

$$\text{cost per IC} = \text{variable cost per IC} + \frac{\text{fixed cost}}{\text{volume}}$$

$$\text{variable cost} = \frac{\text{cost of die} + \text{cost of die test} + \text{cost of packaging}}{\text{final test yield}}$$

Die Cost



$$\text{cost of die} = \frac{\text{cost of wafer}}{\text{dies per wafer} * \text{die yield}}$$


From: <http://www.amd.com>

Yield

$$Y = \frac{\text{No. of good chips per wafer}}{\text{Total number of chips per wafer}} \times 100\%$$


$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{wafer diameter}/2)^2}{\text{die area}} - \frac{\pi \times \text{wafer diameter}}{\sqrt{2} \times \text{die area}}$$




Defects

Yield = 0.25



Yield = 0.76



$$\text{die yield} = \left(1 + \frac{\text{defects per unit area} \times \text{die area}}{\alpha} \right)^{-\alpha}$$

α is approximately in the range 0.5-5

If $\alpha = 3$, die cost = $f(\text{die area})^4$

Summary

- The design process involves traversing the abstraction layers of specification, modeling, architecture, RTL design and physical implementation
- Tests follow the design refinements
- Targets are processors, FPGAs or ASICs
- Automated design flows help manage the complexity
- Optimize for performance, energy and cost